

Introduction to Natural Language Processing Concepts

Nanelle Tuguinay
College of Computer Studies and
Engineering
Jose Rizal University
Pasig City, Philippines
nanelle.tuguinay@my.jru.edu

Abstract—This lab reflection documents the foundational concepts of Natural Language Processing (NLP), focusing on the pipeline required to transform unstructured human text into structured, machine readable representations. As part of this exercise, I completed the Microsoft Learn module "Introduction to natural language processing concepts" (a 7-unit learning path covering tokenization, statistical text analysis, and semantic language models). The report synthesizes key theoretical insights such as tokenization trade-offs, vectorization techniques, and contextual embeddings and examines their practical application in text analytics pipelines and capstone level machine learning models.

Keywords—*Natural Language Processing (NLP), Text Analytics, Tokenization, TF-IDF (Term Frequency-Inverse Document Frequency), Word Embeddings, Context Vectors, Semantic Analysis, Machine Learning, Capstone Applications*

I. INTRODUCTION

Natural Language Processing (NLP) serves as a critical bridge between human communication and computational intelligence, enabling software applications to process, analyze, and infer semantic meaning from unstructured text [1]. Because machine learning algorithms operate strictly on mathematical vectors rather than raw text strings, understanding how to clean, tokenize, and numerically encode language is fundamental to building reliable NLP applications.

Prior to completing the Microsoft Learn training module [2] and reviewing enterprise NLP overviews [1], my baseline understanding of NLP was centered primarily on high-level machine learning concepts and standard supervised classification techniques. I understood that computers required numerical data to process text, but lacked a granular view of how specific preprocessing choices such as subword tokenization, term weighting, or context-aware embeddings directly influence model performance, syntactic parsing, and semantic understanding.

II. EXPERIENTIAL & PRACTICAL INSIGHTS

During the hands on text analytics exercise in Unit 5 (*Exercise - Explore text analytics*), I observed firsthand how raw text processing directly impacts qualitative model outputs:

- **Handling Noise:** Standardizing text through lowercasing and punctuation removal is essential for eliminating statistical redundancy. Without these steps, the same semantic word is falsely

treated as separate vocabulary items due to superficial formatting variations.

- **Frequency Dynamics:** Relying purely on raw word frequency yields poor analytical results because high-frequency functional words dominate the corpus without offering actual thematic context.
- **Feature Filtering:** Filtering out common stop words while preserving core domain terminology dramatically improves the signal-to-noise ratio in text models, allowing meaningful terms to drive the underlying analytics.

III. CORE LEARNING OUTCOMES

Tokenization Trade-offs

Tokenization breaks down unstructured text into discrete processing units (tokens).

- **Word-level tokenization** preserves clear semantic boundaries but struggles heavily with out-of-vocabulary (OOV) terms and creates excessively large dictionary sizes.
- **Character-level tokenization** eliminates OOV issues completely and keeps the vocabulary minimal, but loses semantic context and produces much longer sequences.
- **Subword tokenization** strikes a balance by capturing common sub-words and root structures, maintaining a manageable vocabulary while effectively handling rare or complex word variations.

TF-IDF vs. Word Embeddings

- **TF-IDF (Term Frequency-Inverse Document Frequency):** A sparse, frequency-based statistical vectorizer. It measures how important a word is relative to a document within a broader corpus. However, it lacks semantic awareness (treating words like "car" and "automobile" as entirely independent, unrelated features) and ignores word order completely.
- **Word Embeddings:** Dense, continuous vector representations that project words into a high-dimensional continuous vector space. Words with similar meanings or usages are positioned closer together, allowing models to infer conceptual relationships rather than relying on exact word matches.

Context Vectors

- While basic embeddings assign a fixed static vector to a word regardless of how it is used, context vectors dynamically adapt representation based on surrounding words. This resolves issues of polysemy ensuring that a word like "bank" receives distinct vector representations depending on whether the surrounding context refers to financial institutions or riverbanks.

IV. COURSE APPLICATION & IMPACT

Completing this foundational module provides a strong framework for my ongoing coursework and practical IT developments:

- Upcoming Course Milestones & Labs: Understanding the mechanics of tokenization, normalization, and vector representations provides the technical ground necessary for building custom NLP pipelines, feature extractors, and sequence models in upcoming lab exercises.
- Capstone Application (ProcureSmart): For my capstone project developing a machine learning-based procurement estimation system these concepts are directly applicable. Unstructured text inputs (such as line-item descriptions, vendor notes, equipment details, or procurement requisitions) require careful normalization and vectorization to be processed accurately by predictive models. Applying semantic embeddings

over pure keyword matching will ensure the system can effectively categorize procurement items and match related historical entries even when terminology varies across department requests.

V. CONCLUSION & FUTURE WORK

This exercise reinforced that successful text analytics depends heavily on choosing the right vectorization strategy for the underlying task. While classical statistical methods like TF-IDF remain efficient for basic text classification and keyword search, semantic context models and context vectors are essential for complex language understanding.

Future Study Areas:

- Exploring Transformer-based architectures and attention mechanisms.
- Fine-tuning large language models (LLMs) for specialized tasks.
- Implementing modern vector databases to perform efficient semantic similarity search and Retrieval-Augmented Generation (RAG).

References

- [1] Oracle ASEAN, "What is Natural Language Processing (NLP)?," Oracle ASEAN, [Online].
- [2] Microsoft, "Introduction to natural language processing concepts," Microsoft Learn, [Online Module]